AI IN CLOUD

THE DEFINITIVE STUDY GUIDE

AWS · Google Cloud · Microsoft Azure

Amazon Bedrock · Amazon SageMaker · Vertex AI · Gemini Enterprise Agent Platform
Microsoft Foundry · Copilot Studio · Azure OpenAI Service

Comprehensive Reference for Cloud AI Architecture, Security, Operations & Strategy

Version 2.0 — May 2026

Table of Contents

Part 1: Foundations — Cloud AI Landscape in 2026

1.1 What Is Cloud AI?

Cloud AI refers to artificial intelligence services delivered as managed platform offerings by hyperscale cloud providers. Rather than procuring GPU hardware, installing drivers, configuring inference servers, and managing model lifecycles internally, organizations consume AI capabilities through APIs, managed runtimes, and visual consoles hosted and operated by AWS, Google Cloud, and Microsoft Azure. These services abstract away the infrastructure layer and expose AI functionality as composable building blocks: foundation model inference, fine-tuning pipelines, agentic orchestration, retrieval-augmented generation (RAG), guardrails, evaluation frameworks, and MLOps toolchains.

The cloud AI market in 2026 has matured significantly from the early API-wrapper era of 2023–2024. Today's platforms are full lifecycle systems that span model selection, customization, deployment, monitoring, governance, and multi-agent orchestration. The three dominant ecosystems each bring distinct philosophies: AWS emphasizes model choice and serverless simplicity through Bedrock; Google Cloud unifies everything under the Gemini Enterprise Agent Platform (the evolution of Vertex AI); and Microsoft converges on Microsoft Foundry as a PaaS for agentic AI with deep Microsoft 365 integration.

1.2 Why Cloud AI Over Self-Hosted?

The decision between cloud-managed AI services and self-hosted infrastructure is one of the most consequential architectural choices an organization makes. The following analysis covers the key trade-off dimensions:

Dimension	Cloud AI (Managed)	Self-Hosted
Time to Value	Minutes to hours. API key, SDK call, production endpoint.	Weeks to months. Procure GPUs, install CUDA, configure model serving.
Capital Expenditure	Zero upfront. Pay-per-token or provisioned throughput.	Significant. NVIDIA H100/B200 GPUs cost $25K–$40K each; clusters run into millions.
Operational Overhead	Zero infrastructure management. Provider handles scaling, patching, failover.	Full-time SRE/MLOps team required. Driver updates, thermal management, model versioning.
Model Diversity	Access to 100–200+ models from multiple providers through a single API.	Limited to what you can download, license, and serve yourself.
Security & Compliance	Enterprise-grade: IAM, VPC/PrivateLink, encryption, SOC2, HIPAA, FedRAMP.	Full control but full responsibility. Every compliance certificate is your burden.
Scalability	Elastic. Burst from 0 to 10,000 concurrent requests automatically.	Capacity-bound. Over-provision (waste money) or under-provision (drop requests).
Data Sovereignty	Regions available globally; some services support air-gapped/on-prem deployment.	Complete control over data location and movement.
Customization Depth	Fine-tuning, distillation, RAG, guardrails. Full training possible on SageMaker.	Unlimited. Modify model architecture, training loops, serving infrastructure.
Cost at Scale	Can become expensive at very high token volumes (millions of daily requests).	Lower marginal cost once infrastructure is amortized.

Dimension	Cloud AI (Managed)	Self-Hosted
Vendor Lock-in	Moderate. Unified APIs reduce but don't eliminate switching costs.	None. But you bear all the complexity.

When cloud AI wins: Rapid prototyping, variable workloads, multi-model experimentation, compliance-heavy industries (healthcare, finance), teams without dedicated ML infrastructure engineers, and organizations that need to move fast without building an AI platform team.

When self-hosted wins: Extremely high-volume steady-state inference (millions of requests/day on a single model), air-gapped environments with no cloud connectivity, deep model customization requiring architecture changes, and organizations with existing GPU clusters and MLOps maturity.

The hybrid pattern (most common in 2026): Organizations use cloud AI for exploration, agent building, and variable workloads while running performance-critical, high-volume inference on dedicated infrastructure (often SageMaker HyperPod or GKE with TPUs). SageMaker Unified Studio and Microsoft Foundry both explicitly support this hybrid model.

1.3 The Three Pillars of Cloud AI Architecture

Every cloud AI platform, regardless of provider, is built around three architectural pillars:

- **Model Layer:** The foundation models themselves—proprietary (GPT-5, Claude, Gemini), open-weight (Llama 4, Mistral, Gemma, DeepSeek), and custom (fine-tuned or trained from scratch). This layer includes model catalogs, evaluation tools, and deployment options.
- **Orchestration Layer:** The middleware that connects models to data, tools, and business logic—agents, RAG pipelines, knowledge bases, guardrails, memory systems, and multi-agent workflows. This is where the "agentic AI" revolution of 2025–2026 lives.
- **Operations Layer:** The infrastructure for running AI in production—compute (GPUs, TPUs, Inferentia), networking (VPCs, PrivateLink), security (IAM, encryption, audit logging), observability (tracing, metrics, evaluations), and governance (policies, cost controls, compliance).

Part 2: AWS — Amazon Bedrock, SageMaker & AgentCore

2.1 Amazon Bedrock — The Serverless Foundation Model Platform

2.1.1 What It Is

Amazon Bedrock is AWS's fully managed, serverless service that provides API access to a curated selection of high-performing foundation models from Amazon and leading third-party AI companies through a unified, consistent interface. As of May 2026, Bedrock powers generative AI for over 100,000 organizations worldwide. It operates on a Model-as-a-Service (MaaS) paradigm: you consume models through API calls without provisioning, managing, or scaling any infrastructure.

2.1.2 How It Works

Bedrock's architecture is built around the concept of a GenAI Application Stack where infrastructure complexity is hidden behind high-level abstractions:

- **Unified API (Converse API):** A single, standardized API for multi-turn dialogues across all models. Switch between Claude, Llama, Mistral, or Nova without changing application code.
- **Inference Engine (Mantle):** A new distributed inference engine purpose-built for large-scale model serving. Mantle simplifies model onboarding, provides highly performant serverless inference with quality-of-service controls, and unlocks higher default customer quotas through automated capacity management.
- **Knowledge Bases:** Managed RAG pipelines that automatically ingest, chunk, embed, and index your documents in vector stores (Amazon OpenSearch Serverless, Aurora PostgreSQL, Pinecone). Queries are automatically augmented with retrieved context.
- **Agents:** Autonomous reasoning engines that decompose tasks, select tools, execute actions (via Lambda functions or API calls), and return structured responses. Agents support Action Groups (custom business logic), Code Interpretation, and episodic memory.
- **Guardrails:** Content filtering, PII detection, denied topic blocking, and grounding checks. Apply to any model, including imported custom models.

2.1.3 Available Models (May 2026)

Bedrock now provides nearly 100 serverless models. The April 2026 expansion added 18 fully managed open-weight models, the largest single expansion to date:

Provider	Key Models	Strengths
Amazon	Nova Pro, Nova Lite, Nova Micro, Titan Text, Titan Embeddings, Titan Image Generator	Cost-optimized, native AWS integration, multimodal
Anthropic	Claude Opus 4.7, Claude Sonnet 4.6, Claude Haiku 4.5	Reasoning, coding, 1M token context, agentic workflows
Meta	Llama 4 Maverick, Llama 4 Scout, Llama 3.3	Open-weight, multilingual, customizable

Provider	Key Models	Strengths
Mistral AI	Mistral Large 3, Magistral Small 1.2, Ministral 3 (3B/8B/14B), Voxtral Mini/Small	Multilingual, tool use, edge-optimized variants
OpenAI (Preview)	GPT-5.5, GPT-5.4, GPT OSS 120B/20B, Codex	Frontier reasoning, coding agent, open-weight options
Google	Gemma 3 (various sizes)	Open-weight, efficient, research-friendly
NVIDIA	Nemotron Nano 9B/12B/30B v2, Nemotron Super 120B	High compute efficiency, agentic systems
Cohere	Command R+, Embed v4, Rerank	Enterprise search, RAG, multilingual
AI21 Labs	Jamba 2 family	Long context, enterprise text
Stability AI	Stable Diffusion 3.5, SDXL	Image generation
Others	DeepSeek, Qwen, MiniMax, Moonshot AI, Luma AI, TwelveLabs, Writer	Specialized use cases

2.1.4 Security Architecture

Bedrock inherits AWS's enterprise security model and adds AI-specific protections:

- **Identity & Access Management (IAM):** Fine-grained policies controlling who can invoke which models, create agents, or access knowledge bases. Service-linked roles for automated operations.
- **Network Isolation:** AWS PrivateLink endpoints ensure model traffic never traverses the public internet. VPC configurations isolate Bedrock resources within your network boundary.
- **Encryption:** Data encrypted at rest (AWS KMS with customer-managed keys) and in transit (TLS 1.2+). Model customization data is encrypted with dedicated keys.
- **Audit Logging:** CloudTrail captures every API call. Model invocation logging records prompts and responses (opt-in) for compliance and debugging.
- **Data Privacy:** Your data is not used to train base models. Fine-tuning data remains within your AWS account. Model providers cannot access your prompts or completions.
- **Guardrails:** Configurable content filters, PII redaction (with customizable regex/entity types), denied topic enforcement, and contextual grounding checks. Applied at the API layer, independent of the underlying model.
- **Compliance:** SOC 1/2/3, ISO 27001, HIPAA eligible, PCI DSS, FedRAMP (GovCloud), GDPR-ready. Available in AWS GovCloud (US) for government workloads.

2.1.5 Deployment & Consumption Patterns

Bedrock offers three deployment types:

- **On-Demand:** Pay-per-token with no commitment. Ideal for variable workloads, experimentation, and applications with unpredictable traffic. No provisioning required.
- **Provisioned Throughput:** Purchase model units for guaranteed, consistent performance. Ideal for production workloads with predictable traffic patterns and latency SLA requirements. Billed hourly regardless of usage.

- **Batch Inference:** Submit large datasets for asynchronous processing at up to 50% cost reduction. Ideal for document processing, bulk content generation, and evaluation pipelines.

2.1.6 AgentCore — The Agent Infrastructure Layer

Launched in late 2025 and significantly expanded in April 2026, Amazon Bedrock AgentCore is the infrastructure backbone for deploying production AI agents on AWS. AgentCore provides:

- **Managed Harness (Preview):** Define an agent by specifying a model, system prompt, and tools, then run it immediately with no orchestration code. The harness manages the full agent loop: reasoning, tool selection, action execution, and response streaming. Each session gets its own microVM with filesystem and shell access.
- **AgentCore CLI:** Command-line interface for creating, testing, and deploying agents from your terminal. Supports exporting harness orchestration as Strands-based code for full customization.
- **Identity & Access:** Every agent gets its own IAM identity, logging each action for audit. Agents run in your environment with all inference on Bedrock.
- **Evaluations:** 13 built-in quality dimensions (helpfulness, accuracy, tool selection, hallucination detection) for continuous agent quality monitoring.
- **Episodic Memory:** Agents learn from past interactions, improving performance over time without retraining.
- **Policy Controls:** Define agent boundaries using natural language that converts to Cedar policy language.

2.1.7 Amazon Bedrock Managed Agents, Powered by OpenAI

Announced April 28, 2026, this offering combines frontier OpenAI models with the OpenAI agent harness on AWS infrastructure. The harness is engineered for faster execution, sharper reasoning, and reliable steering of long-running tasks. Customers authenticate with AWS credentials, all inference runs through Bedrock, and usage counts toward existing AWS cloud commitments. This represents a significant expansion of the AWS-OpenAI partnership.

2.2 Amazon SageMaker AI — The Full ML Lifecycle Platform

2.2.1 What It Is

Amazon SageMaker AI (formerly Amazon SageMaker) is a fully managed service designed for building, training, and deploying machine learning models at scale. While Bedrock operates at the application layer (consuming pre-trained models via API), SageMaker operates at the infrastructure layer (building, training, and customizing models from the ground up). SageMaker is the destination for organizations that need to train proprietary models, perform massive-scale distributed training, or require specialized inference hardware.

2.2.2 Key Capabilities

- **SageMaker Unified Studio:** A single development environment that integrates Bedrock, SageMaker, data analytics, and ML workflows under one project structure. Access both Bedrock IDE and SageMaker Studio from the same workspace.
- **SageMaker HyperPod:** Resilient, multi-node training across thousands of accelerators. Automatically repairs hardware failures to ensure trillion-parameter training jobs don't crash. Supports Slurm and Amazon EKS orchestrators.
- **JumpStart:** A model hub with hundreds of open-source models plus support for custom models. Deploy pre-trained models for inference or fine-tuning with one click.
- **Training & Tuning:** Managed services for training, including distributed training, automatic mixed precision, spot training (up to 90% cost savings), and hyperparameter optimization.
- **Inference Options:** Real-time endpoints (dedicated instances), serverless inference (auto-scaling), batch transform (bulk processing), and asynchronous inference (queued workloads). AWS Inferentia3 chips provide cost-optimized inference for specific models.
- **MLOps:** Pipelines (CI/CD for ML), Model Registry (versioning and approval workflows), Model Monitor (data drift, bias detection, feature attribution), and MLflow integration for experiment tracking.
- **Canvas:** No-code ML for business analysts. Point-and-click interface for predictions without writing code.

2.2.3 When to Use Bedrock vs. SageMaker

Decision Factor	Choose Bedrock	Choose SageMaker
Primary Goal	Build GenAI applications (chatbots, RAG, content gen)	Build/train custom ML models from scratch
Team Expertise	Application developers with API experience	Data scientists with ML/framework expertise
Customization	Fine-tuning, RAG, guardrails, model distillation	Architecture changes, custom training loops, RLHF
Infrastructure	Serverless, zero management	Managed instances, GPU/TPU provisioning
Cost Model	Pay-per-token (variable)	Instance-hours + training compute (fixed + variable)
Agent Building	Bedrock Agents + AgentCore (built-in)	Bring your own framework (LangChain, Strands)

Decision Factor	Choose Bedrock	Choose SageMaker
Time to Market	Hours to days	Weeks to months
Scale Pattern	Variable, bursty workloads	Steady-state, high-throughput inference

The convergence story: As of 2026, the boundaries are blurring. Bedrock now supports reinforcement learning fine-tuning (RFT), and SageMaker has added serverless customization with RLVR and RLAIF. SageMaker Unified Studio lets teams use both from the same workspace. The recommendation: start with Bedrock for speed, graduate to SageMaker when you need deeper customization or cost optimization at scale.

Part 3: Google Cloud — Vertex AI & Gemini Enterprise Agent Platform

3.1 Vertex AI — The Unified AI Platform

3.1.1 What It Is

Vertex AI is Google Cloud's unified, open platform for building, deploying, and scaling generative AI and machine learning models and AI applications. In 2026, Vertex AI is transitioning to become part of the Gemini Enterprise Agent Platform, announced at Google Cloud Next 2026. This rebranding reflects a shift from model-centric tooling to an agent-first platform that bundles model selection, agent building, orchestration, governance, and runtime into a single offering.

Vertex AI provides access to the Model Garden with over 200 models, including Google's Gemini family, Anthropic Claude (Opus, Sonnet, Haiku as first-class citizens), Meta Llama, open-source Gemma, and dozens of partner and community models. It supports both generative AI workflows and traditional ML/MLOps through a comprehensive toolset.

3.1.2 Key Components

- **Model Garden:** A curated catalog of 200+ models from Google, partners, and the open-source community. Browse, evaluate, deploy, and fine-tune models from a single interface.
- **Vertex AI Studio:** A visual interface for prompt design, model experimentation, and rapid prototyping. Supports text, image, audio, and video modalities.
- **Agent Development Kit (ADK):** An open-source, graph-based framework for building and orchestrating multi-agent systems. Define agents as networks of sub-agents with clear, reliable logic. ADK is model-agnostic—swap models without rewriting agents.
- **Agent Studio:** A low-code visual builder for designing agents. Export logic directly to ADK for full-code customization when needed.
- **Agent Engine (Runtime):** A managed, serverless runtime for deploying agents at scale. Handles compute provisioning, scaling, and lifecycle management.
- **Sessions & Memory Bank (GA):** Persistent memory for agents across sessions. Automatic extraction, consolidation, and retrieval of context. Billed at $0.25 per 1,000 events/memories.
- **Vertex AI Search:** Enterprise search with vector, full-text, and semantic re-ranking. Powers RAG pipelines with grounded, citation-backed responses.
- **Model Armor:** Runtime defense against prompt injection, data exfiltration, and other emergent threats.
- **Gen AI Evaluation Service:** Evaluate and compare model and agent performance. Supports partner models (Anthropic, Llama) alongside Google models.

3.1.3 Available Models (May 2026)

Model Family	Key Variants	Capabilities
Gemini 3.1 (Preview)	Pro, Flash-Lite	Reasoning-first, 1M context, adaptive thinking, integrated grounding, agentic workflows

Model Family	Key Variants	Capabilities
Gemini 2.5	Pro, Flash, Flash-Lite, Flash Image	Complex reasoning, 1M context, controllable thinking budgets, image generation
Gemini 2.0	Flash, Flash-Lite	Multimodal, general-purpose, cost-effective (legacy, existing customers only from Mar 2026)
Gemini Live API	2.5 Flash Native Audio	Real-time bidirectional streaming, native audio, affective dialog, proactive audio
Veo	3.1, 3.1 Lite, 3.1 Fast	Video generation at various quality/cost trade-offs
Imagen / Gemini Image	3.1 Flash Image	Image generation, conversational editing, multi-image fusion, character consistency
Lyria 3	Pro, Clip (Preview)	Audio generation (up to 184 seconds)
Gemma 3	Various sizes	Open-weight, 140+ languages, 128K context, text and image input
Partner Models	Claude Opus 4.7, Sonnet 4.6, Haiku 4.5; Llama 4; DeepSeek V3.2; Kimi K2; GLM 5; MiniMax M2	Full spectrum of third-party capabilities

3.1.4 Security & Governance

- **IAM & RBAC:** Google Cloud IAM with fine-grained roles for model access, agent management, and data operations. Agent identities for managing access and authentication.
- **VPC Service Controls:** Define security perimeters around Vertex AI resources. Prevent data exfiltration through service boundaries.
- **Customer-Managed Encryption Keys (CMEK):** Encrypt data with keys you control via Cloud KMS. As of March 2026, CMEK is available at the service level for Azure AI Search indexes.
- **Data Residency:** Regional deployment options. Google Distributed Cloud (GDC) supports fully disconnected, on-premises deployment for sovereign and air-gapped environments.
- **Model Armor:** Proactive runtime inspection against prompt injection, jailbreaking, and data exfiltration. Configurable safety filters with content moderation policies.
- **Audit Logging:** Cloud Audit Logs capture all administrative and data access events. Agent observability includes sessions, traces, logs, and events.
- **Compliance:** SOC 1/2/3, ISO 27001/27017/27018, HIPAA, PCI DSS, FedRAMP, and 50+ regional certifications.

3.1.5 Deployment & Consumption

- **Pay-as-you-go:** Per-token pricing for model inference. No flat subscription fee. Foundation model costs are typically the largest line item.
- **Provisioned Throughput:** Purchase Generative AI Scale Units (GSUs) for guaranteed throughput on specific models. Supports global endpoints for traffic distribution.

- **Agent Engine Runtime:** Billed on compute consumption (vCPU-hours starting at $0.0864/vCPU-hour). Sessions and memories billed separately.
- **Express Mode (Free Tier):** Free tier for Agent Engine Runtime for experimentation and development.
- **Infrastructure:** TPU 8t (training, 3x compute vs. previous gen) and TPU 8i (inference, 80% better perf-per-dollar for agentic workloads). GPU options via A3 and A4 VM families.

3.1.6 Gemini Enterprise Agent Platform

Announced at Google Cloud Next 2026, this is the evolution of Vertex AI into a comprehensive agent platform. Beyond the existing model selection and building capabilities, it adds:

- **Agent Integration:** Cloud API Registry for managing MCP servers and tools. Managed remote MCP servers for databases. Tools for Data Agents with high-accuracy text-to-SQL.
- **Agent DevOps:** CI/CD for agents. Version control, testing, staged rollouts, and rollback.
- **Agent Orchestration:** Multi-agent workflows with clear delegation patterns. A2A (Agent-to-Agent) protocol support.
- **Agent Security:** Model Armor at the runtime layer. Agent identity management. Policy enforcement for agent actions.
- **Sovereign Neocloud:** Google Distributed Cloud for on-premises and air-gapped deployments. Large multimodal models on local NVIDIA hardware without cloud connectivity.

Part 4: Microsoft Azure — Foundry, Copilot Studio & Azure OpenAI

4.1 Microsoft Foundry — The Unified AI PaaS

4.1.1 What It Is

Microsoft Foundry (formerly Azure AI Foundry, formerly Azure AI Studio) is Microsoft's unified platform-as-a-service for enterprise AI operations, model building, and application development. It combines production-grade infrastructure with developer-friendly interfaces, unifying agents, models, and tools under a single management grouping with built-in tracing, monitoring, evaluations, and customizable enterprise setup. Foundry is used by developers at more than 80,000 enterprises and digital natives, including 80% of Fortune 500 companies.

The platform evolution reflects Microsoft's strategic shift from isolated experimentation tools to a fleet-wide AI management system designed for the agentic era. The February 2026 update elevated Foundry from experimentation tool to production-capable solution with multi-agent orchestration, MCP support, hosted agents, and sovereign local deployment.

4.1.2 Key Components

- **Foundry Models:** A catalog of 11,000+ models from Microsoft, OpenAI, Anthropic Claude, Meta, Mistral AI, DeepSeek, xAI (Grok), Cohere, HuggingFace, NVIDIA, Moonshot AI, and more. Serverless pay-as-you-go or managed compute deployment.
- **Foundry Agent Service (GA, March 2026):** The next-gen agent runtime, built on the OpenAI Responses API and wire-compatible with OpenAI agents. Supports models from DeepSeek, xAI, Meta, LangChain, LangGraph, and more. Enterprise security, private networking, Entra RBAC, full tracing, and evaluation built-in.
- **Foundry IQ:** Knowledge integration layer that grounds agent responses in enterprise or web content with citation-backed answers. Powered by Azure AI Search, connects agents to multiple data sources with built-in user access permissions.
- **Foundry Tools:** Over 1,400 connectors to enterprise systems (SAP, Salesforce, Dynamics 365). Pre-built tools for OCR, translation, speech, and object detection. MCP support for custom APIs.
- **Memory in Agent Service (Preview):** Managed long-term memory with automatic extraction, consolidation, and retrieval across agent sessions.
- **Voice Live:** Fully managed real-time speech-to-speech runtime. Collapses STT-to-LLM-to-TTS pipeline into a single API with voice activity detection, noise suppression, echo cancellation, and barge-in support.
- **Foundry MCP Server (Preview):** Cloud-hosted MCP at mcp.ai.azure.com. Zero local process management, Entra auth included. Connect from VS Code, Visual Studio, or the portal.

4.1.3 Available Models (May 2026)

Provider	Key Models	Notes
Azure OpenAI	GPT-5-chat, GPT-5.5 Instant, GPT-5.2 (GA), GPT-4.1 series, GPT OSS 120B, Codex Max (GA)	Frontier reasoning, coding, multimodal
Microsoft Phi	Phi family (60M+ downloads)	Small language models, low-cost, on-device capable
Anthropic	Claude Opus 4.6, Sonnet 4.6, Haiku 4.5	Deep reasoning, coding, safety
Meta	Llama 4 Maverick, Llama 4 Scout	Open-weight, multilingual
Mistral AI	Mistral Large, Ministral series	Multilingual, efficient
DeepSeek	V3.2, V3.2 Speciale	128K context, sparse attention, reasoning
xAI	Grok-4, Grok-code-fast-1	Enterprise reasoning, coding, data extraction
Moonshot AI	Kimi K2.5 (Preview)	256K context, deep reasoning, tool orchestration
Cohere	Rerank v4.0 (Fast/Pro)	Enterprise search re-ranking
NVIDIA	Nemotron family	Fine-tunable, edge-deployable

4.1.4 Deployment Types

- **Standard (Global):** Routes traffic globally for higher throughput. Best for applications without strict data residency requirements.
- **Data Zone Standard:** Traffic stays within a defined geographic zone. Addresses data sovereignty requirements while maintaining managed infrastructure.
- **Provisioned (Global):** Purchase provisioned throughput units for guaranteed, consistent performance. Available globally.
- **Hosted Agents:** Foundry provisions compute, registers endpoints, provides a production URL. Currently limited to North Central US (as of March 2026).
- **Bring-Your-Own-Runtime:** Deploy conventional App Service, Container App, or AKS workloads while calling Foundry's /openai/v1/ endpoints directly. Separates application runtime from Foundry infrastructure.
- **On-Premises (Sovereign):** Large multimodal models on local NVIDIA hardware without cloud connectivity. Supports air-gapped environments via Azure Local and Azure Stack HCI.

4.1.5 Security Architecture

- **Microsoft Entra ID (Azure AD):** Centralized identity with RBAC. Agent Service supports full Entra authentication for MCP and A2A protocols.
- **Azure Key Vault:** Secrets management for API keys, connection strings, and certificates used by agents and models.
- **Azure Private Link:** Private connectivity to Foundry endpoints. No public internet exposure.
- **Microsoft Defender Integration:** AI-specific threat detection. Rich signals management for security, performance, cost, and governance.

- **Azure Policy:** Governance guardrails applied at the subscription or management group level. Enforce model restrictions, region constraints, and configuration standards.
- **Customer-Managed Keys (CMK):** Encryption at rest with keys you control. Service-level CMK for Azure AI Search (new March 2026).
- **PII NextGen:** Conversational and document PII detection with advanced configuration (March 2026).
- **Compliance:** SOC 1/2/3, ISO 27001, HIPAA, PCI DSS, FedRAMP, GDPR, and 50+ global certifications.

4.2 Microsoft Copilot Studio — The Low-Code Agent Builder

4.2.1 What It Is

Microsoft Copilot Studio is a SaaS agent platform that enables organizations to build AI agents and agentic workflows to transform business processes without deep coding expertise. It offers a graphical, low-code authoring canvas with drag-and-drop topic design, natural language flow configuration, and seamless publishing to Microsoft 365, Teams, BizChat, web, and custom channels. It is distinct from Microsoft Foundry in that it targets citizen developers and business analysts rather than professional developers.

4.2.2 Key Capabilities (2026)

- **Multi-Agent Orchestration (GA, April 2026):** Build connected, governed multi-agent systems. Agent-to-agent calling via agent nodes in workflows.
- **Model Choice:** Use GPT models (ChatGPT-5 GA globally), Claude Sonnet 4.5/4.6 and Opus (GA globally), and other external models as the primary AI for agents.
- **Generative Answers:** Point agents at internal/external knowledge sources. AI automatically creates conversational responses without manually authored topics.
- **Generative Actions:** AI dynamically selects and chains plugins in real-time based on user intent.
- **Agent Flows & Workflows:** Standalone automations or agent-triggered tools. Natural language or visual designer. Trigger manually, by events, or on schedule.
- **Content Moderation:** Configurable harmful content sensitivity. Adjustable for regulated industries (healthcare, insurance, law enforcement).
- **Prompt Assistant:** GPT-powered prompt drafting suggestions for faster, more effective prompt engineering.
- **Publishing:** One-click deployment to Microsoft 365 Copilot, Teams, BizChat, web chat, or containerized deployments.
- **Analytics:** Session transcripts, CSAT values, question/reaction filtering, downloadable CSV reports.

4.2.3 Copilot Studio vs. Microsoft Foundry

Dimension	Copilot Studio	Microsoft Foundry
Target User	Citizen developers, business analysts, IT admins	Professional developers, ML engineers, platform teams
Authoring	Low-code visual canvas, natural language	Code-first (Python, JS, .NET, Java SDKs), YAML, CLI
Model Access	GPT, Claude, selected external models	11,000+ models from 20+ providers
Agent Complexity	Conversational agents, FAQ bots, simple workflows	Complex multi-agent systems, custom harnesses, agentic RAG
Integration	Microsoft 365, Power Platform, 1,400+ connectors	Azure services, any REST API, MCP, A2A protocols
Deployment	Microsoft-managed SaaS	Your Azure subscription with full infrastructure control

Dimension	Copilot Studio	Microsoft Foundry
Pricing	Per-message licensing (included in some M365 plans)	Pay-as-you-go per service consumed
Governance	Built-in DLP, compliance via Power Platform admin center	Azure Policy, Entra ID, Defender, full Azure governance stack

Part 5: Cross-Cloud Comparison & Decision Framework

5.1 Platform Comparison Matrix

Capability	AWS (Bedrock/SageMaker)	Google Cloud (Vertex AI)	Azure (Foundry)
Model Catalog Size	~100 serverless + JumpStart (hundreds more)	200+ in Model Garden	11,000+ in Foundry Models
First-Party Models	Amazon Nova, Titan	Gemini family, Gemma	GPT (via OpenAI), Phi
Agent Framework	Bedrock Agents + AgentCore + Strands SDK	ADK (open-source) + Agent Engine	Foundry Agent Service (Responses API) + MCP/A2A
Low-Code Agent Builder	Bedrock IDE (in Unified Studio)	Agent Studio	Copilot Studio
RAG / Knowledge	Knowledge Bases (managed)	Vertex AI Search + Foundry IQ equiv.	Foundry IQ + Azure AI Search
Guardrails	Bedrock Guardrails (content, PII, grounding)	Model Armor	Azure AI Content Safety + PII NextGen
Memory	Episodic memory (agents)	Sessions + Memory Bank (GA)	Memory in Agent Service (Preview)
Voice/Audio	Amazon Connect + Lex	Gemini Live API (native audio, GA)	Voice Live (managed STT/LLM/TTS)
On-Premises	AWS Outposts	Google Distributed Cloud (GDC)	Azure Local / Stack HCI
MCP Support	Via AgentCore	Cloud API Registry + managed MCP servers	Foundry MCP Server (cloud-hosted)
MLOps Depth	SageMaker (deep: pipelines, monitor, HyperPod)	Vertex AI Training, Workbench, Colab Enterprise	Azure ML (transitioning to Foundry)
Unique Strength	Model choice breadth + serverless simplicity	TPU infrastructure + Gemini native models	Microsoft 365 integration + enterprise reach

5.2 Decision Framework: Which Platform When?

Choosing a cloud AI platform is rarely a binary decision. Most enterprises adopt a primary platform aligned with their existing cloud investment while maintaining the ability to use specific capabilities from others. The following framework helps identify the primary platform:

- **Choose AWS Bedrock when:** Your organization is AWS-native, you need the broadest serverless model selection, you want the fastest path from idea to production API, your workloads are variable/bursty, or you need GovCloud support. Add SageMaker when you need custom model training at scale.
- **Choose Google Cloud Vertex AI when:** You want Gemini models natively (especially for long context, multimodal, and real-time audio), you need TPU infrastructure for

training, your data lives in BigQuery, you want the open-source ADK for agent orchestration, or you need sovereign deployment via GDC.

- **Choose Microsoft Foundry when:** Your organization runs on Microsoft 365 and Azure, you need deep integration with Teams/Outlook/SharePoint, you want the largest model catalog, you need citizen developers to build agents (Copilot Studio), or you require GPT/OpenAI models with enterprise governance.
- **Multi-cloud pattern:** Use Foundry for employee-facing agents integrated with Microsoft 365, Bedrock for customer-facing applications with diverse model needs, and Vertex AI for data-intensive workloads grounded in BigQuery. Each platform’s SDK is independent; multi-cloud is an integration challenge, not a technical barrier.

Part 6: Security Deep-Dive — Securing Cloud AI at Scale

6.1 The AI-Specific Threat Model

Cloud AI introduces threat vectors that don’t exist in traditional application security. A comprehensive security strategy must address:

- **Prompt Injection:** Adversarial inputs that cause models to ignore instructions, leak system prompts, or execute unintended actions. Mitigated by guardrails (Bedrock), Model Armor (Vertex AI), and Content Safety (Azure).
- **Data Exfiltration via Agents:** Agents with tool access could be manipulated into extracting sensitive data through API calls, file reads, or search queries. Mitigated by agent identity policies (AgentCore), IAM scoping, and output filtering.
- **Training Data Poisoning:** Manipulated fine-tuning datasets that embed backdoors or biases into customized models. Mitigated by data validation pipelines, evaluation frameworks, and model testing.
- **Model Theft/Extraction:** Repeated queries designed to reconstruct model weights or distill capabilities. Mitigated by rate limiting, usage monitoring, and anomaly detection.
- **Supply Chain Risk:** Third-party models from the catalog could contain vulnerabilities or unexpected behaviors. Mitigated by model evaluation tools, sandboxed testing, and provenance tracking.
- **Compliance Drift:** AI systems generating content that violates industry regulations (HIPAA patient data, PCI cardholder data, GDPR personal data). Mitigated by guardrails/content filters, PII detection, audit logging, and output monitoring.

6.2 Security Controls by Cloud Provider

Security Control	AWS Bedrock/SageMaker	Google Cloud Vertex AI	Microsoft Foundry
Identity	IAM policies, service roles	Cloud IAM, agent identities	Entra ID, RBAC, managed identities
Network	PrivateLink, VPC endpoints	VPC-SC, Private Google Access, PSC	Private Link, VNet integration
Encryption (Rest)	KMS (CMK supported)	Cloud KMS (CMEK)	Key Vault (CMK), service-level CMK
Encryption (Transit)	TLS 1.2+	TLS 1.2+	TLS 1.2+
Prompt/Output Filtering	Guardrails (content, PII, topics, grounding)	Model Armor (injection, exfiltration)	Content Safety, PII NextGen
Audit	CloudTrail, model invocation logs	Cloud Audit Logs, agent traces	Azure Monitor, Defender, agent traces
Data Sovereignty	Regional endpoints, GovCloud	Regional, GDC (air-gapped)	Regional, Data Zones, Azure Local
Compliance Certs	SOC, ISO, HIPAA, PCI, FedRAMP	SOC, ISO, HIPAA, PCI, FedRAMP, 50+	SOC, ISO, HIPAA, PCI, FedRAMP, 50+

Security Control	AWS Bedrock/SageMaker	Google Cloud Vertex AI	Microsoft Foundry
Agent-Specific	Cedar policies, per-agent IAM	Agent identity IAM, Model Armor runtime	Entra auth for MCP/A2A, Defender AI
Secret Management	Secrets Manager, Parameter Store	Secret Manager	Key Vault

6.3 Security Best Practices for Cloud AI

- **Principle of Least Privilege:** Grant models and agents access only to the specific resources they need. Use separate IAM roles for different agent functions. Never use admin credentials for agent operations.
- **Defense in Depth:** Layer network isolation (PrivateLink/VPC-SC), identity controls (IAM), content filtering (guardrails/Model Armor), output monitoring (audit logs), and anomaly detection. No single control is sufficient.
- **Zero-Trust for Agents:** Treat every agent action as potentially adversarial. Validate inputs, constrain outputs, log everything, and require human approval for high-impact actions (financial transactions, data deletions, external communications).
- **Data Classification:** Classify data before exposing it to AI systems. Sensitive data (PII, PHI, financial records) requires additional guardrails, encryption, and access controls. Never fine-tune models on unclassified data dumps.
- **Continuous Evaluation:** Run automated evaluations against your deployed models and agents. Monitor for hallucination rates, safety violations, and performance degradation. All three clouds provide evaluation frameworks—use them.
- **Incident Response:** Extend your security incident response playbooks to cover AI-specific scenarios: model misbehavior, prompt injection attacks, data leakage through agent actions, and compliance violations in generated content.
- **Supply Chain Hygiene:** Vet third-party models before production deployment. Test with adversarial inputs, evaluate on your specific data, and monitor behavior over time. Track model versions and deprecation schedules.

Part 7: Use Cases, Architecture Patterns & Implementation

7.1 Common Use Cases by Industry

Industry	Use Case	Recommended Platform(s)	Key Services
Financial Services	Fraud detection, compliance review, risk analysis, document intelligence	Bedrock (RAG, Guardrails), Foundry (M365 integration)	Knowledge Bases, Guardrails, PII detection
Healthcare	Clinical documentation, patient engagement, drug interaction analysis	Vertex AI (HIPAA, Med models), Bedrock (GovCloud)	MedGemma, Guardrails, encrypted endpoints
Retail/E-commerce	Product recommendations, customer service bots, inventory optimization	Bedrock (diverse models, low-latency)	Agents, Knowledge Bases, Nova models
Manufacturing	Predictive maintenance, quality inspection, supply chain optimization	SageMaker (custom ML), Vertex AI (vision models)	HyperPod, AutoML, Gemini Vision
Legal	Contract analysis, case research, compliance monitoring	Foundry (document AI), Bedrock (long-context models)	Claude Opus (200K), Foundry IQ, RAG
Education	Tutoring agents, content generation, assessment	Copilot Studio (low-code), Vertex AI (Gemini)	Agent Studio, generative answers
Government	Citizen services, document processing, intelligence analysis	Bedrock (GovCloud, FedRAMP), Foundry (Azure Gov)	GovCloud models, guardrails, audit logging
Software Development	Code generation, review, testing, documentation	Foundry (Codex, GPT), Bedrock (Claude, Codex)	Codex, Claude Opus 4.7, GitHub Copilot

7.2 Architecture Patterns

7.2.1 RAG (Retrieval-Augmented Generation)

The most common cloud AI pattern. Ground model responses in your organization's data to reduce hallucination and provide domain-specific answers.

- **AWS:** Bedrock Knowledge Bases (managed ingest, chunking, embedding, vector store). Supports OpenSearch Serverless, Aurora PostgreSQL, Pinecone, Redis.
- **Google Cloud:** Vertex AI Search with vector, full-text, and hybrid search. RAG Cross Corpus Retrieval (preview) for multi-source queries. BigQuery-native Gemma embeddings.
- **Azure:** Foundry IQ powered by Azure AI Search. Securely ground agents on enterprise and web content with citation-backed answers and built-in user access permissions.

7.2.2 Multi-Agent Systems

Chain specialized agents for complex business processes. Each agent handles a specific domain (document intake, verification, system update, notification) with clear handoff patterns.

- **AWS:** Bedrock Managed Agents with AgentCore. Define agents via managed harness or Strands SDK. Each agent gets its own microVM.
- **Google Cloud:** Agent Development Kit (ADK) with graph-based orchestration. Define agent networks with sub-agent delegation and clear routing logic.
- **Azure:** Foundry Agent Service with A2A (Agent-to-Agent) protocol. Multi-agent workflows built visually in the portal. Copilot Studio agent nodes for cross-agent calling.

7.2.3 Conversational AI with Voice

- **AWS:** Amazon Connect + Amazon Lex + Bedrock. Connect for telephony, Lex for speech recognition, Bedrock for intelligent responses.
- **Google Cloud:** Gemini Live API with native audio (GA). Real-time bidirectional streaming, affective dialog, proactive audio. No separate STT/TTS pipeline needed.
- **Azure:** Voice Live for managed speech-to-speech. Connect directly to Foundry agents—same prompt, tools, and tracing as text-based agents. Copilot Studio voice channel with post-call actions.

7.2.4 Code Generation & Software Engineering

- **AWS:** Codex on Bedrock (preview), Claude Opus 4.7 (1M context, agentic coding). Amazon Q Developer for IDE integration.
- **Google Cloud:** Gemini Code Assist (IDE plugin). Gemini 3.1 Pro optimized for coding. CodeGemma for open-source code completion.
- **Azure:** GitHub Copilot (powered by GPT/Codex). Copilot in Visual Studio 2026 with agent skills, plan-based coding, and multi-file diffs. Codex Max (GA) in Foundry.

Part 8: Operations, Governance & Cost Management

8.1 GenAIOps — The New DevOps

Operating AI systems in production requires a new operational discipline that extends traditional DevOps with AI-specific practices:

- **Model Lifecycle Management:** Track model versions, deprecation schedules, and migration paths. All three clouds deprecate and retire models on published timelines (typically 12–18 months after successor release). Automate migration testing.
- **Evaluation Pipelines:** Continuously evaluate deployed models against quality, safety, and cost metrics. Use provider evaluation tools (Bedrock Evaluations, Gen AI Evaluation Service, Foundry evals) plus custom benchmarks on your data.
- **Observability:** Instrument agents and applications with tracing, logging, and metrics. All three platforms now provide agent-level observability (sessions, traces, events). Integrate with existing APM tools (Datadog, New Relic, Splunk).
- **Cost Monitoring:** Token consumption can be unpredictable. Set budgets, alerts, and quotas. Use prompt caching (all platforms support it) to reduce repeated context costs. Evaluate model routing (use cheaper models for simple queries).
- **Safety Monitoring:** Track guardrail trigger rates, content safety violations, and user feedback. Investigate spikes in blocked content—they may indicate adversarial activity or legitimate use case gaps.

8.2 Cost Optimization Strategies

- **Model Selection:** Match model capability to task complexity. Don't use GPT-5 for intent classification when Phi or Ministral 3B will suffice. Use the evaluation tools to find the cheapest model that meets your quality threshold.
- **Prompt Caching:** Cache large system prompts and repeated context across requests. All platforms support this. Saves 50–90% on input token costs for applications with stable system prompts.
- **Intelligent Prompt Routing:** Route simple queries to small/fast models and complex queries to frontier models. Bedrock offers this natively; Google and Azure support it via agent logic.
- **Batch Processing:** Use batch inference (up to 50% cheaper on Bedrock) for non-real-time workloads like document processing, evaluation runs, and content generation.
- **Provisioned Throughput:** For steady-state workloads, provisioned throughput is cheaper than on-demand at scale. Calculate your break-even point based on daily token volume.
- **Spot/Preemptible Training:** Use spot instances (SageMaker) or preemptible VMs (Vertex AI) for training jobs. Up to 90% savings with automatic checkpointing and restart.

8.3 Governance Framework

- **Model Registry:** Maintain a centralized registry of approved models, versions, and configurations. Document why each model was selected, what data it was fine-tuned on, and what guardrails are applied.

- **Access Controls:** Implement role-based access at every layer: who can deploy models, who can create agents, who can access production endpoints, who can modify guardrails. Use existing identity providers (Entra ID, Google IAM, AWS IAM).
- **Policy Enforcement:** Define organizational AI policies covering acceptable use, data classification, output review requirements, and escalation procedures. Enforce programmatically via Azure Policy, AWS Service Control Policies, or GCP Organization Policies.
- **Responsible AI:** Implement bias detection, fairness evaluation, and toxicity monitoring. Use provider tools (SageMaker Clarify, Responsible AI Toolkit on Vertex AI) and supplement with domain-specific evaluations.
- **Documentation & Audit Trail:** Maintain model cards documenting capabilities, limitations, and intended use. Keep audit logs of all model invocations, agent actions, and human overrides. These are increasingly required by regulations (EU AI Act, NIST AI RMF).

Part 9: Emerging Trends & Future Direction

9.1 The Agentic Era

2026 marks the definitive shift from “AI as API” to “AI as Agent.” All three clouds have invested heavily in agent infrastructure: managed runtimes, memory systems, tool protocols (MCP, A2A), multi-agent orchestration, and agent-specific governance. The key trends:

- **Agent Interoperability:** MCP (Model Context Protocol) is emerging as the standard for connecting agents to tools and data sources. All three clouds now support MCP—Azure’s Foundry MCP Server, Google’s managed MCP servers for databases, and AWS’s AgentCore MCP integration.
- **Agent-to-Agent (A2A):** Protocols for agents to call other agents across organizational boundaries. Microsoft and Google are leading with A2A protocol support.
- **Agent Identity & Governance:** Every agent gets its own identity, permissions, and audit trail. This is critical for compliance and accountability in multi-agent systems.

9.2 Model Convergence & Commoditization

The model landscape is converging: all platforms now offer models from most major providers. OpenAI models are available on AWS Bedrock (preview), Microsoft Foundry, and (indirectly) via API. Anthropic models are available on all three clouds. Meta, Mistral, and DeepSeek are ubiquitous. This commoditization shifts the competitive differentiator from model access to orchestration quality, developer experience, and ecosystem integration.

9.3 Sovereign & Air-Gapped AI

Data sovereignty requirements are driving all three clouds to offer on-premises AI deployment: AWS Outposts, Google Distributed Cloud, and Azure Local/Stack HCI. The February 2026 expansion of Azure’s local deployment capabilities—from small language models to large multimodal models on NVIDIA hardware—signals that sovereign AI is no longer a compromise on capability.

9.4 The SDK Consolidation

All three clouds are consolidating their AI SDKs. Microsoft shipped unified 2.0.0 SDKs (Python, JS/TS, Java, .NET) in March 2026, eliminating separate packages for agents, inference, evaluations, and memory. AWS is converging Bedrock and SageMaker under Unified Studio. Google is unifying under the Gemini Enterprise Agent Platform umbrella. The developer experience is becoming simpler even as the underlying capabilities grow more complex.

Part 10: Quick Reference & Glossary

10.1 Key Terminology

Term	Definition
Foundation Model (FM)	A large-scale model pre-trained on broad data, designed to be adapted for downstream tasks.
RAG	Retrieval-Augmented Generation. Grounding model responses in retrieved documents to reduce hallucination.
Agent	An AI system that reasons, plans, selects tools, and executes actions autonomously to achieve goals.
MCP	Model Context Protocol. A standard for connecting AI models/agents to external tools and data sources.
A2A	Agent-to-Agent. A protocol for agents to invoke other agents across boundaries.
Guardrails	Content filtering, PII detection, and safety mechanisms applied to model inputs and outputs.
Fine-Tuning	Adapting a pre-trained model to a specific task or domain using additional training on specialized data.
Provisioned Throughput	Pre-purchased compute capacity guaranteeing consistent performance for model inference.
Model Garden / Catalog	A curated collection of foundation models available for deployment through a cloud platform.
Vector Store	A database optimized for storing and querying high-dimensional embeddings for similarity search.
Inference	The process of running input data through a trained model to generate predictions or outputs.
MLOps / GenAIOps	Operational practices for deploying, monitoring, and maintaining ML/AI systems in production.
SLM	Small Language Model. A compact model (1B–30B parameters) optimized for specific tasks with lower compute requirements.
Context Window	The maximum number of tokens a model can process in a single request (prompt + response).
Token	The basic unit of text processing. Roughly 3/4 of a word in English. Pricing is typically per-token.

10.2 Service Name Mapping

Capability	AWS	Google Cloud	Microsoft Azure
FM API Platform	Amazon Bedrock	Vertex AI / Gemini Enterprise Agent Platform	Microsoft Foundry (Foundry Models)
Custom ML Platform	Amazon SageMaker AI	Vertex AI Training / Workbench	Azure Machine Learning (migrating to Foundry)

Capability	AWS	Google Cloud	Microsoft Azure
Agent Builder (Code)	Bedrock AgentCore + Strands SDK	Agent Development Kit (ADK)	Foundry Agent Service + Foundry SDK
Agent Builder (Low-Code)	Bedrock IDE	Agent Studio	Copilot Studio
Knowledge / RAG	Bedrock Knowledge Bases	Vertex AI Search	Foundry IQ + Azure AI Search
Safety / Guardrails	Bedrock Guardrails	Model Armor	Azure AI Content Safety
Model Catalog	Bedrock Model Choice + JumpStart	Model Garden (200+)	Foundry Models (11,000+)
Unified Workspace	SageMaker Unified Studio	Google Cloud Console / Colab Enterprise	Foundry Portal + VS Code Extension
Vector Database	OpenSearch Serverless, Aurora pgvector	Vector Search 2.0, AlloyDB	Azure AI Search, Cosmos DB (vector)
Sovereign / On-Prem	AWS Outposts	Google Distributed Cloud (GDC)	Azure Local / Stack HCI

10.3 Key URLs & Documentation

- **AWS Bedrock:** docs.aws.amazon.com/bedrock — Console: console.aws.amazon.com/bedrock
- **AWS SageMaker:** docs.aws.amazon.com/sagemaker — Console: console.aws.amazon.com/sagemaker
- **Google Vertex AI:** cloud.google.com/vertex-ai/docs — Console: console.cloud.google.com/vertex-ai
- **Microsoft Foundry:** learn.microsoft.com/azure/foundry — Portal: ai.azure.com
- **Copilot Studio:** learn.microsoft.com/microsoft-copilot-studio — Portal: copilotstudio.microsoft.com
- **Azure OpenAI:** learn.microsoft.com/azure/ai-services/openai

---

End of Study Guide — Version 2.0 — May 2026

This document should be reviewed and updated quarterly as cloud AI services evolve rapidly.